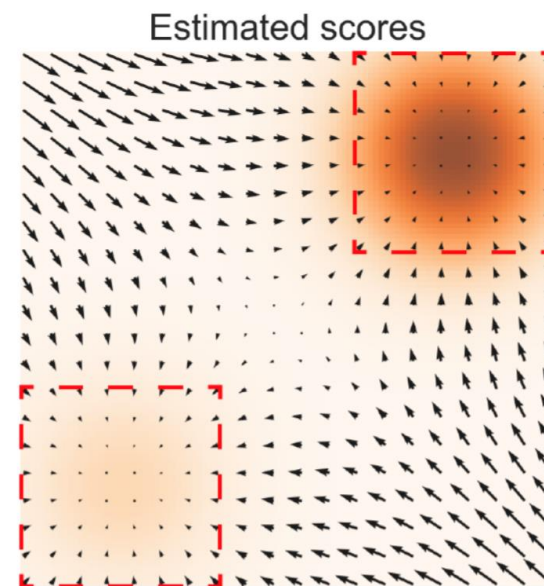
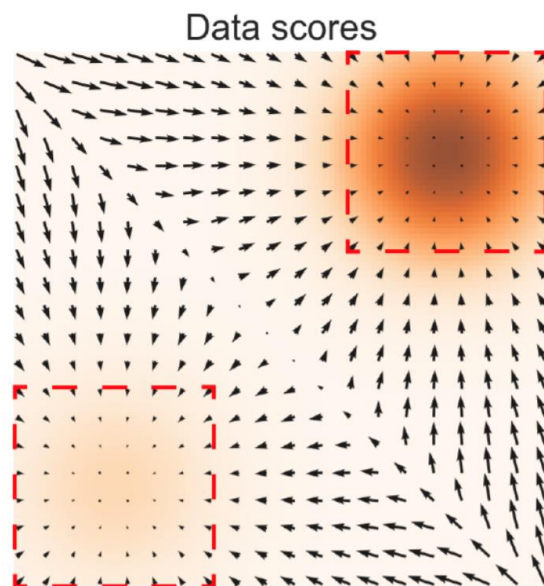
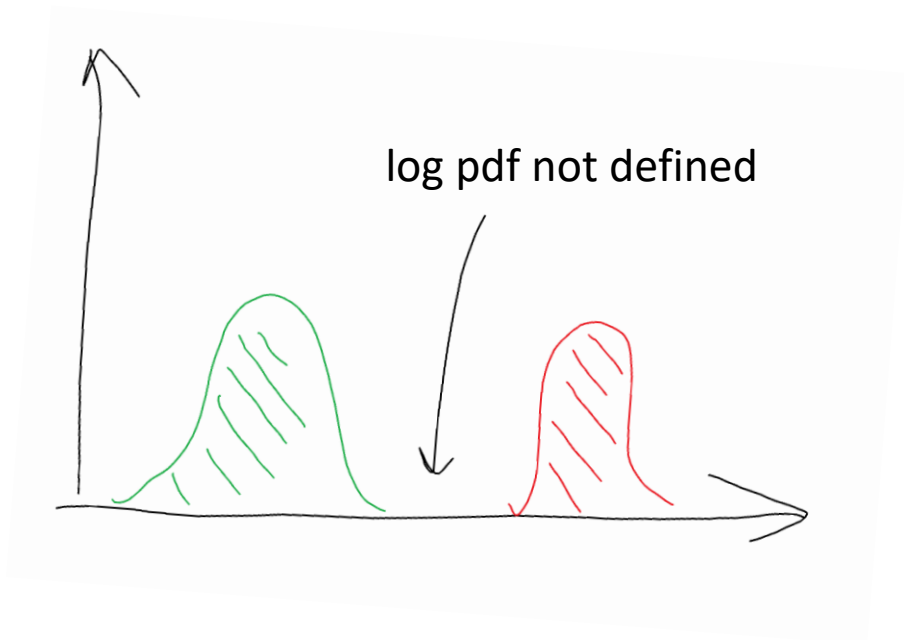


Diffusion Models

Anthony Baryshnikov

As score matching models

- Score matching estimates the gradient of log prob in data space and generates samples by some variation of gradient ascent.
- It is good but has some downsides.
- Extra backpropagations to estimate trace of score gradient.
- Manifold hypothesis.
- Low data density regions have bad score estimates.
- It's hard to transition between disconnected supports.



As score matching models

- Let's perturb data with various levels of Gaussian noise and train a noise conditional network to estimate score.
- High noise fills low density regions and gives a common support.
- Low noise is almost indistinguishable from true data.
- Sample using annealed Langevin dynamics (ALS).
- Our loss becomes:

$$\ell(\boldsymbol{\theta}; \sigma) \triangleq \frac{1}{2} \mathbb{E}_{p_{\text{data}}(\mathbf{x})} \mathbb{E}_{\tilde{\mathbf{x}} \sim \mathcal{N}(\mathbf{x}, \sigma^2 I)} \left[\left\| \mathbf{s}_{\boldsymbol{\theta}}(\tilde{\mathbf{x}}, \sigma) + \frac{\tilde{\mathbf{x}} - \mathbf{x}}{\sigma^2} \right\|_2^2 \right] \quad \mathcal{L}(\boldsymbol{\theta}; \{\sigma_i\}_{i=1}^L) \triangleq \frac{1}{L} \sum_{i=1}^L \lambda(\sigma_i) \ell(\boldsymbol{\theta}; \sigma_i)$$

Algorithm 1 Annealed Langevin dynamics.

Require: $\{\sigma_i\}_{i=1}^L, \epsilon, T$.

1: Initialize $\tilde{\mathbf{x}}_0$

2: **for** $i \leftarrow 1$ to L **do**

3: $\alpha_i \leftarrow \epsilon \cdot \sigma_i^2 / \sigma_L^2$ $\triangleright \alpha_i$ is the step size.

4: **for** $t \leftarrow 1$ to T **do**

5: Draw $\mathbf{z}_t \sim \mathcal{N}(0, I)$

6: $\tilde{\mathbf{x}}_t \leftarrow \tilde{\mathbf{x}}_{t-1} + \frac{\alpha_i}{2} \mathbf{s}_\theta(\tilde{\mathbf{x}}_{t-1}, \sigma_i) + \sqrt{\alpha_i} \mathbf{z}_t$

7: **end for**

8: $\tilde{\mathbf{x}}_0 \leftarrow \tilde{\mathbf{x}}_T$

9: **end for**

return $\tilde{\mathbf{x}}_T$

Too many hyperparameters

- How can we come up with good noise levels?
- What about sampling step size?
- And the number of steps?

Noise levels

- Smallest noise level has to be indistinguishable.
- Transition probabilities decay exponentially.
- Choose largest noise level at least as large as the maximum Euclidean distance between all pairs of training data.
- Samples have to cover high density regions of previous noise level.
- Choose a geometric progression with common ratio dependent on data dimensionality.

ALS parameters

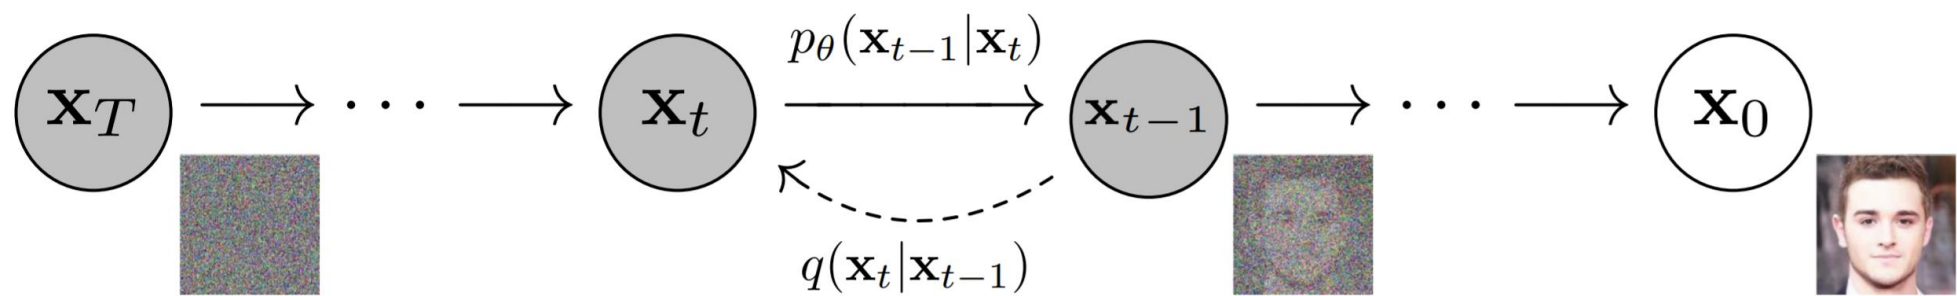
- We want sampling variance to be as close to noise level as possible.
- Can be computed analytically for one data point.
- Choose T as large as possible and optimize the step size making variance ratio as close to 1 as possible.

Other tricks

- Hard to condition on noise level (is it really?).
- Make an unconditional score estimation network and divide its output by noise std.
- Samples are empirically very unstable and exhibit artifacts such as common color shift.
- Apply EMA over model weights.

As nonequilibrium thermodynamics

- Let's gradually perturb our data with small noise.
- Reverse diffusion process has the same functional form.
- We have to predict mean and variance.
- Train by maximizing variational lower bound.
- Generate by gradually denoising samples from stationary distribution.



$$L \geq \int d\mathbf{x}^{(0\cdots T)} q\left(\mathbf{x}^{(0\cdots T)}\right) \cdot$$

$$\log \left[p\left(\mathbf{x}^{(T)}\right) \prod_{t=1}^T \frac{p\left(\mathbf{x}^{(t-1)}|\mathbf{x}^{(t)}\right)}{q\left(\mathbf{x}^{(t)}|\mathbf{x}^{(t-1)}\right)} \right]$$

$$K = - \sum_{t=2}^T \int d\mathbf{x}^{(0)} d\mathbf{x}^{(t)} q\left(\mathbf{x}^{(0)}, \mathbf{x}^{(t)}\right) \cdot$$

$$D_{KL}\left(q\left(\mathbf{x}^{(t-1)}|\mathbf{x}^{(t)}, \mathbf{x}^{(0)}\right) \Big\| \Big| p\left(\mathbf{x}^{(t-1)}|\mathbf{x}^{(t)}\right)\right)$$

$$+ H_q\left(\mathbf{X}^{(T)}|\mathbf{X}^{(0)}\right) - H_q\left(\mathbf{X}^{(1)}|\mathbf{X}^{(0)}\right) - H_p\left(\mathbf{X}^{(T)}\right) \cdot$$

As nonequilibrium thermodynamics

- We're now working with Gaussian noise only.
- Let's set variance to a time dependent constant.

Loss reparameterization

- Let's rewrite our loss.
- We can remove the factor between estimated noise difference norm to obtain a simplified objective.
- Score matching objective and variance lower bound maximization are very similar.

$$\mathbb{E}_{\mathbf{x}_0, \epsilon} \left[\frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1 - \bar{\alpha}_t)} \left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2 \right]$$

Algorithm 1 Training

```
1: repeat  
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$   
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$   
4:    $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
5:   Take gradient descent step on  
        $\nabla_{\theta} \left\| \boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t) \right\|^2$   
6: until converged
```

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
2: for  $t = T, \dots, 1$  do  
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$   
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$   
5: end for  
6: return  $\mathbf{x}_0$ 
```

How to obtain exact log likelihoods?

- Let's set the last term of reversed process to a discrete decoder.
- We can now estimate the conditional probability by calculating an integral.




$$p_{\theta}(\mathbf{x}_0|\mathbf{x}_1) = \prod_{i=1}^D \int_{\delta_{-}(x_0^i)}^{\delta_{+}(x_0^i)} \mathcal{N}(x; \mu_{\theta}^i(\mathbf{x}_1, 1), \sigma_1^2) dx$$
$$\delta_{+}(x) = \begin{cases} \infty & \text{if } x = 1 \\ x + \frac{1}{255} & \text{if } x < 1 \end{cases} \quad \delta_{-}(x) = \begin{cases} -\infty & \text{if } x = -1 \\ x - \frac{1}{255} & \text{if } x > -1 \end{cases}$$

Different objectives

- Simplified objective makes low noise level loss relatively more important and improves sample quality.
- Predicting noise gives similar results to estimating mean of Gaussian.

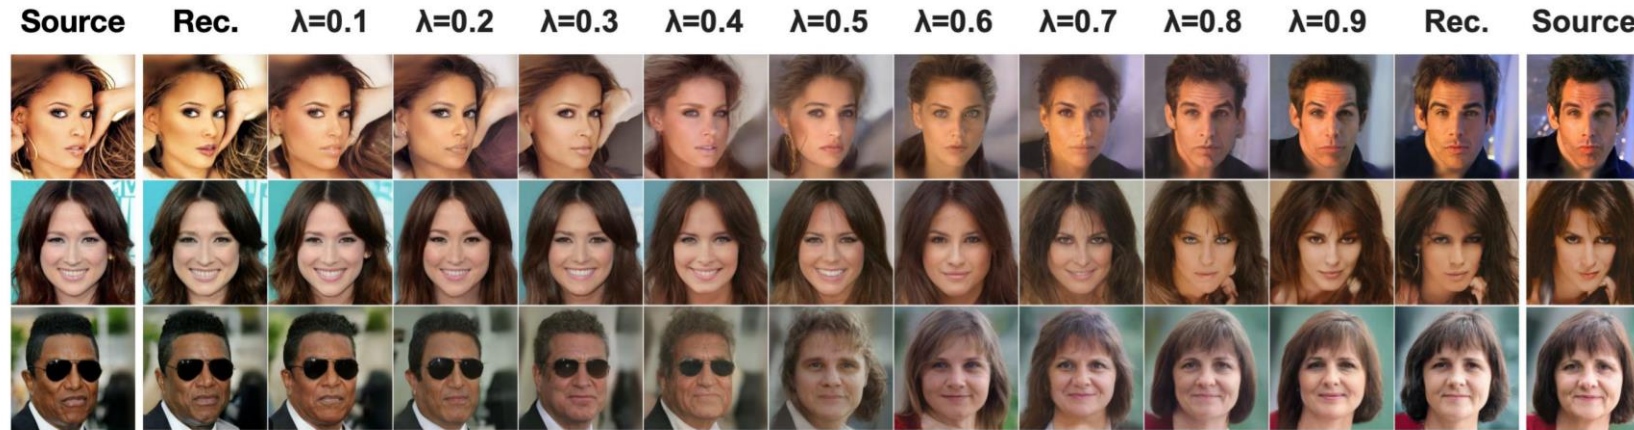
Lossy compression

- KL divergence corresponds to rate (?).
- RMSE corresponds to distortion.
- Let's plot them.
- Turns out that the majority of our codelength encodes impeccable details, which is not optimal.
- I'm not sure if I got this right.



Extra details

- Diffusion process that masks first T pixels corresponds to an autoregressive model.
- Interpolating in latent space works particularly well.

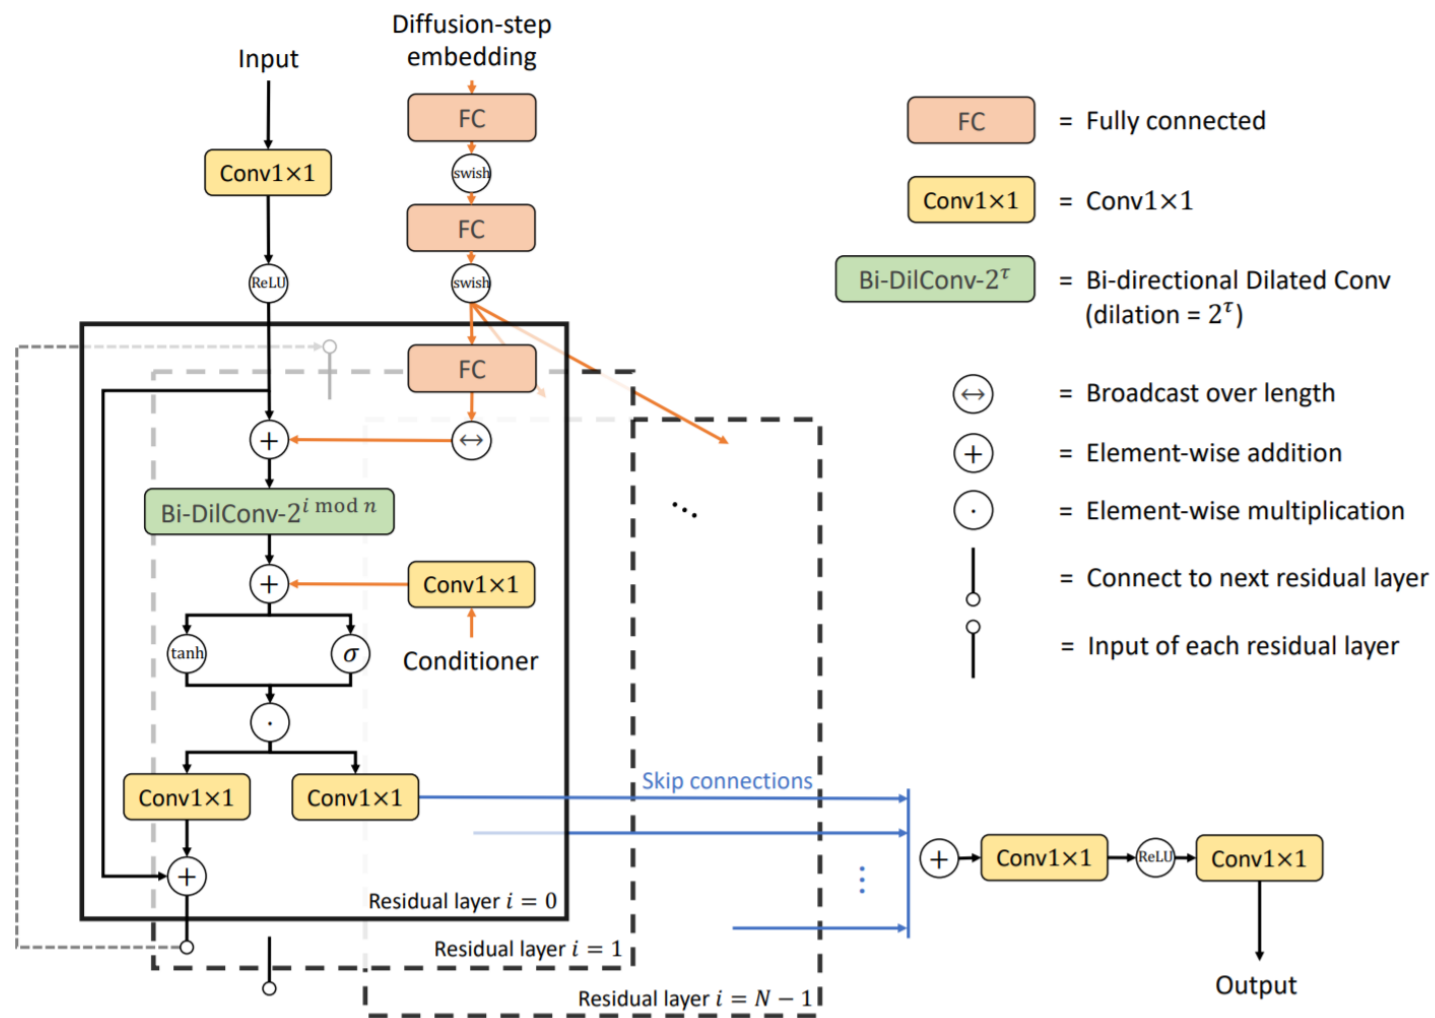


Advantages in audio generation

- Diffusion models avoid mode/posterior collapse.
- Non-autoregressive which means faster parallel synthesis.
- Very flexible architecture.
- Does not require auxiliary losses (e.g. Mel spectrogram loss).
- Provides intuitive speed/quality tradeoff by varying number of denoising steps.

DiffWave

- Bidirectional dilated convolutions.
- N residual layers divided into m blocks.
- Dilation is doubled at each layer within each block.
- Big receptive fields because of multiple denoising steps.
- Positional embeddings of timesteps from transformers.
- Let's upsample local conditioning to the same length.
- We can then add both local and global conditioning as bias terms after dilated convolutions by using 1×1 conv.
- We can use faster noise schedules at inference to increase speed.



DiffWave results

- Performs well at neural vocoding.
- Better MOS than WaveNet at 6.91M vs 4.57M parameters.
- Real-time generation but still much slower than flow based models (1.1-5.6x vs 40+x).
- Completely destroys everybody at unconditional and class-conditional generation.
- Zero-shot speech denoising and latent space interpolation is also available.

Table 1: The model hyperparameters, model footprint, and 5-scale Mean Opinion Score (MOS) with 95% confidence intervals for WaveNet, ClariNet, WaveFlow, WaveGlow and the proposed DiffWave on the **neural vocoding** task. \uparrow means the number is the higher the better, and \downarrow means the number is the lower the better.

Model	T	T_{infer}	layers	res. channels	#param(\downarrow)	MOS(\uparrow)
WaveNet	—	—	30	128	4.57M	4.43 \pm 0.10
ClariNet	—	—	60	64	2.17M	4.27 \pm 0.09
WaveGlow	—	—	96	256	87.88M	4.33 \pm 0.12
WaveFlow	—	—	64	64	5.91M	4.30 \pm 0.11
WaveFlow	—	—	64	128	22.25M	4.40 \pm 0.07
DiffWave _{BASE}	20	20	30	64	2.64M	4.31 \pm 0.09
DiffWave _{BASE}	40	40	30	64	2.64M	4.35 \pm 0.10
DiffWave _{BASE}	50	50	30	64	2.64M	4.38 \pm 0.08
DiffWave _{LARGE}	200	200	30	128	6.91M	4.44 \pm 0.07
DiffWave _{BASE} (Fast)	50	6	30	64	2.64M	4.37 \pm 0.07
DiffWave _{LARGE} (Fast)	200	6	30	128	6.91M	4.42 \pm 0.09
Ground-truth	—	—	—	—	—	4.52 \pm 0.06

Table 2: The automatic evaluation metrics (FID, IS, mIS, AM, and NDB/ K), and 5-scale MOS with 95% confidence intervals for WaveNet, WaveGAN, and DiffWave on the **unconditional** generation task. \uparrow means the number is the higher the better, and \downarrow means the number is the lower the better.

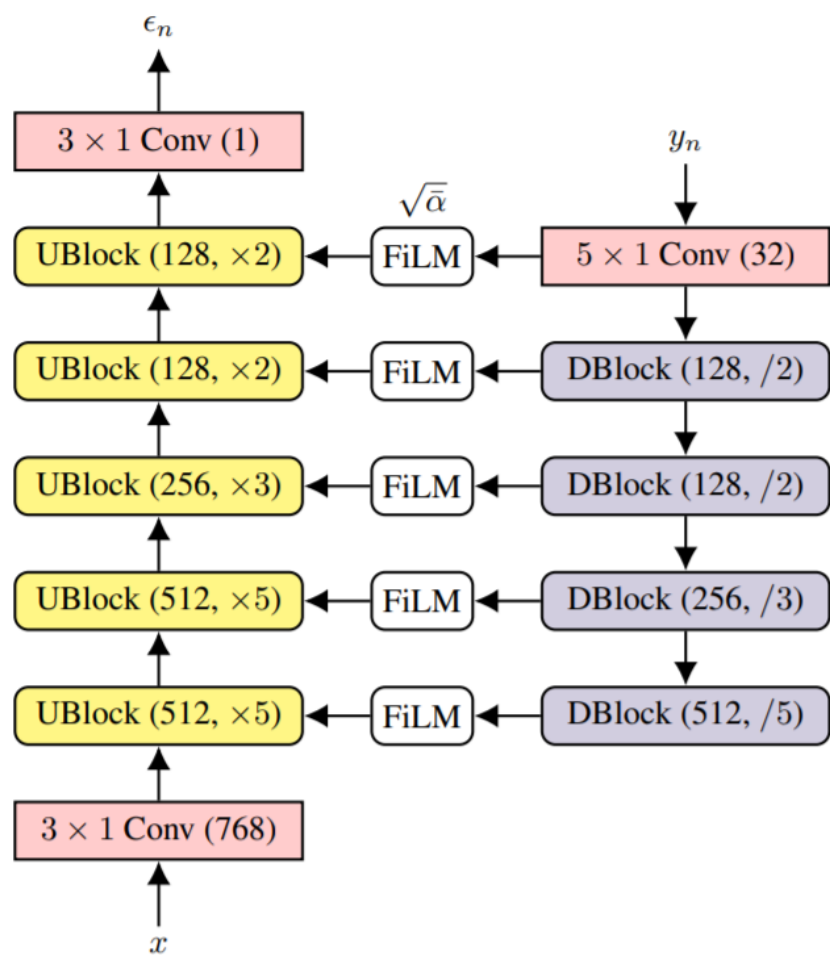
Model	FID(\downarrow)	IS(\uparrow)	mIS(\uparrow)	AM(\downarrow)	NDB/ K (\downarrow)	MOS(\uparrow)
WaveNet-128	3.279	2.54	7.6	1.368	0.86	1.34 ± 0.29
WaveNet-256	2.947	2.84	10.0	1.260	0.86	1.43 ± 0.30
WaveGAN	1.349	4.53	36.6	0.796	0.78	2.03 ± 0.33
DiffWave	1.287	5.30	59.4	0.636	0.74	3.39 ± 0.32
Trainset	0.000	8.48	281.4	0.164	0.00	—
Testset	0.011	8.47	275.2	0.166	0.10	3.72 ± 0.28

Table 3: The automatic evaluation metrics (Accuracy, FID-class, IS, mIS), and 5-scale MOS with 95% confidence intervals for WaveNet and DiffWave on the **class-conditional** generation task.

Model	Accuracy(\uparrow)	FID-class(\downarrow)	IS(\uparrow)	mIS(\uparrow)	MOS(\uparrow)
WaveNet-128	56.20%	7.876 ± 2.469	3.29	15.8	1.46 ± 0.30
WaveNet-256	60.70%	6.954 ± 2.114	3.46	18.9	1.58 ± 0.36
DiffWave	91.20%	1.113 ± 0.569	6.63	117.4	3.50 ± 0.31
DiffWave (deep & thin)	94.00%	0.932 ± 0.450	6.92	133.8	3.44 ± 0.36
Trainset	99.06%	0.000 ± 0.000	8.48	281.4	—
Testset	98.76%	0.044 ± 0.016	8.47	275.2	3.72 ± 0.28

WaveGrad

- Let's use diffusion models in TTS.
- Network similar to a feature pyramid.
- Uses spatial feature-wise linear modulation for conditioning.
- Proposes conditioning on the fraction of true signal instead of on the timestamp which provides better generalization between noise schedules.
- Authors also suggest Fibonacci and manual noise schedules.



WaveGrad results

- Large model with 1000 iterations achieves a MOS of 4.51 (4.58 GT).
- Base model with 6 iterations achieves a MOS of 4.41 with good real-time factors (0.2 on NVIDIA V100 and 1.5 on CPU).

Model	MOS (\uparrow)
WaveRNN	4.49 ± 0.04
Parallel WaveGAN	3.92 ± 0.05
MelGAN	3.95 ± 0.06
Multi-band MelGAN	4.10 ± 0.05
GAN-TTS	4.34 ± 0.04
WaveGrad	
Base (6 iterations, continuous noise levels)	4.41 ± 0.03
Base (1,000 iterations, discrete indices)	4.47 ± 0.04
Large (1,000 iterations, discrete indices)	4.51 ± 0.04
Ground Truth	4.58 ± 0.05

Iterations (schedule)	LS-MSE (\downarrow)	MCD (\downarrow)	FFE (\downarrow)	MOS (\uparrow)
WaveGrad conditioned on discrete index				
25 (Fibonacci)	283	3.93	3.3%	3.86 ± 0.05
50 (Linear (1×10^{-4} , 0.05))	181	3.13	3.1%	4.42 ± 0.04
1,000 (Linear (1×10^{-4} , 0.005))	116	2.85	3.2%	4.47 ± 0.04
WaveGrad conditioned on continuous noise level				
6 (Manual)	217	3.38	2.8%	4.41 ± 0.04
25 (Fibonacci)	185	3.33	2.8%	4.44 ± 0.04
50 (Linear (1×10^{-4} , 0.05))	177	3.23	2.7%	4.43 ± 0.04
1,000 (Linear (1×10^{-4} , 0.005))	106	2.85	3.0%	4.46 ± 0.03

Questions?